

# Producer-Consumer problem

## 1- Solution pseudocode:

The problem describes two processes, the producer and the consumer, who share a common, fixed-size buffer used as a queue.

The producer's job is to generate data and put it into the buffer. At the same time, the consumer consumes the data (i.e., by removing it from the buffer), one piece at a time.

The problem is to make sure that the producer won't try to add data into the buffer if it's full and that the consumer won't try to remove data from an empty buffer.

The solution for the producer is to either go to sleep or discard data if the buffer is full. The next time the consumer removes an item from the buffer, it notifies the producer, who starts to fill the buffer again.

In the same way, the consumer can go to sleep if it finds the buffer to be empty. The next time the producer puts data into the buffer, it wakes up the sleeping consumer

## We will solve this problem using semaphore

we are going to use Binary Semaphore and Counting Semaphore

*Binary Semaphore:* In Binary Semaphore, only two processes can compete to enter its CRITICAL SECTION at any point in time, apart from this the condition of mutual exclusion is also preserved.

*Counting Semaphore:* In counting semaphore, more than two processes can compete to enter its CRITICAL SECTION at any point of time apart from this the condition of mutual exclusion is also preserved

The structure of the producer process:

```
while (true) {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
}
```


- First the producer waits until there is at least one empty slot.
- Then it decreases the empty semaphore because, there will now be one less empty slot, since the producer is going to insert data in one of those slots.
- Then, it acquires lock on the buffer, so that the consumer cannot access the buffer until the producer completes its operation.
- After performing the insert operation, the lock is released and the value of full is incremented because the producer has just filled a slot in the buffer

The structure of the consumer process:

```
while (true) {  
    wait(full);  
    wait(mutex);  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /* consume the item in next consumed */  
    ...  
}
```

- The consumer waits until there is at least one full slot in the buffer.
- Then it decrements the full semaphore because the number of occupied slots will be decreased by one, after the consumer completes its operation.
- After that, the consumer acquires lock on the buffer.
- Following that, the consumer completes the removal operation so that the data from one of the full slots is removed.
- Then, the consumer releases the lock.
- Finally, the empty semaphore is incremented by 1, because the consumer has just removed data from an occupied slot, thus making it empty

## 2- Examples of Deadlock:



A deadlock occurs when both the producer and consumer processes *are waiting for the other to do something*, resulting in a standstill. For example, if the buffer is full and the producer process is waiting for the consumer process to remove an item from the buffer, and the consumer process is waiting for the producer process to add an item to the buffer, a deadlock will occur.

To prevent deadlocks, it is necessary to have a mechanism in place to ensure that either the producer or the consumer can proceed, even if the other is not ready. This can be done using **semaphores** or other **synchronization tools**.

Example:

```
class Buffer {  
    private final BlockingQueue<Integer> queue;  
    private final Object lock1 = new Object();  
    private final Object lock2 = new Object();  
}
```

```

public Buffer(int size) {
    queue = new ArrayBlockingQueue<>(size);
}

public void add(int item) throws InterruptedException {
    synchronized (lock1) {
        while (queue.size() == queue.remainingCapacity()) {
            lock1.wait();
        }
        queue.put(item);
        lock2.notifyAll();
    }
}

public int remove() throws InterruptedException {
    synchronized (lock2) {
        while (queue.size() == 0) {
            lock2.wait();
        }
        int item = queue.take();
        lock1.notifyAll();
        return item;
    }
}
}

```

In this example, the Buffer class uses two separate locks, lock1 and lock2, to **synchronize access** to the shared buffer. The add method acquires **lock1** and waits for the buffer to have space, while the remove method acquires **lock2** and waits for the buffer to have an item.

However, if the producer thread is holding lock1 and waiting for the buffer to have space, and the consumer thread is holding lock2 and waiting for the buffer to have an item, a deadlock will occur. The producer thread will not release lock1 until an item is removed from the buffer, but the consumer thread will not remove an item from the buffer until lock1 is released. As a result, both threads will be stuck waiting indefinitely, **resulting in a deadlock**.

To solve the deadlock problem in the previous code, you can use a **single lock** to **synchronize access** to the shared buffer. Here is an example of how this can be done

```
class Buffer {
    private final BlockingQueue<Integer> queue;
    private final Object lock = new Object();

    public Buffer(int size) {
        queue = new ArrayBlockingQueue<>(size);
    }

    public void add(int item) throws InterruptedException {
        synchronized (lock) {
            while (queue.size() == queue.remainingCapacity()) {
                lock.wait();
            }
            queue.put(item);
            lock.notifyAll();
        }
    }

    public int remove() throws InterruptedException {
        synchronized (lock) {
            while (queue.size() == 0) {
                lock.wait();
            }
            int item = queue.take();
            lock.notifyAll();
            return item;
        }
    }
}
```

In this example, the Buffer class uses a single lock, lock, to synchronize access to the shared buffer. Both the add and remove methods acquire the lock and wait for the buffer to have space or an item, respectively.

Using a single lock ensures that the producer and consumer threads are correctly synchronized and reduces the risk of a deadlock. However, it is still important to carefully design the synchronization mechanisms and ensure that the producer and consumer threads are correctly coordinated to prevent other types of synchronization issues

Another example of a deadlock:

```
class Account {
    double balance;

    void withdraw(double amount){
        balance -= amount;
    }

    void deposit(double amount){
        balance += amount;
    }

    void transfer(Account from, Account to, double amount){
        sync(from);
        sync(to);

        from.withdraw(amount);
        to.deposit(amount);

        release(to);
        release(from);
    }
}
```

Solve the problem:

```
public static void transfer(Account from, Account to, double amount){
    Account first = from;
    Account second = to;
    if (first.compareTo(second) < 0) {
        // Swap them
        first = to;
        second = from;
    }
    synchronized(first){
        synchronized(second){
            from.withdraw(amount);
            to.deposit(amount);
        }
    }
}
```

### 3- Examples of starvation:

It is possible for the producer-consumer problem to result in a starvation problem, where one of the producer or consumer processes is **unable to make progress** because it is constantly being blocked by the other process.

For example, if the producer process is **generating items faster** than the consumer process can consume them, the buffer may fill up and the producer process may be blocked, unable to add more items to the buffer. This could result in the consumer process being starved of items to consume.

Similarly, if the consumer process is consuming items faster than the producer process can generate them, the buffer may empty out and the consumer process may be blocked, waiting for the producer process to add more items to the buffer. This could result in the producer process being starved of an opportunity to produce items.

To prevent starvation in the producer-consumer problem, it is important to carefully design the synchronization mechanisms and ensure that both the producer and consumer processes have a fair opportunity to access the shared buffer. This can be done using **semaphores or other synchronization tools**

Example:

```
class Buffer {
    private final BlockingQueue<Integer> queue;
    public Buffer(int size) {
        queue = new ArrayBlockingQueue<>(size);
    }
    public void add(int item) throws InterruptedException {
        while (true) {
            if (queue.offer(item)) {
                break;
            }
            Thread.sleep(1);
        }
    }
    public int remove() throws InterruptedException {
        while (true) {
            if (queue.size() > 0) {
                return queue.poll();
            }
            Thread.sleep(1);
        }
    }
}
```



In this example, the **Producer** class generates items at a rate of one per second, while the **Consumer** class consumes items at a rate of one per two seconds. This means that the producer generates items faster than the consumer can consume them, and the buffer may fill up.

If the buffer fills up, the add method in the Buffer class will block indefinitely, waiting for an item to be removed from the buffer. This means that the producer process may be starved of an opportunity to produce items, and it may be unable to make progress.

To prevent this starvation problem, it is important to carefully design the synchronization mechanisms and ensure that both the producer and consumer processes have a fair opportunity to access the shared buffer. One option is to use a lock or a semaphore to synchronize access to the buffer, so that the producer and consumer processes take turns accessing it. Another option is to use a blocking queue, which handles synchronization internally and does not require the use of separate locks

To solve the starvation problem in the previous example using a semaphore, you can modify the **Buffer** class to use a semaphore to synchronize access to the shared buffer. Here is an example of how this can be done:

```
class Buffer {
    private final BlockingQueue<Integer> queue;
    private final Semaphore semaphore;
    public Buffer(int size) {
        queue = new ArrayBlockingQueue<>(size);
        semaphore = new Semaphore(1);
    }

    public void add(int item) throws InterruptedException {
        semaphore.acquire();
        queue.put(item);
        semaphore.release();
    }

    public int remove() throws InterruptedException {
        semaphore.acquire();
        int item = queue.take();
        semaphore.release();
        return item;
    }
}
```

## 4- Real World Application

One real-world example of the producer-consumer problem is A car assembly line

**A car assembly line:** In a car assembly line, the producer might be a robot that adds car parts to the assembly line, while the consumer is a person who installs the parts on the car. The common buffer in this case might be the assembly line itself, which holds the car parts as they are produced.

The producer process (robot) would produce car parts and place them on the assembly line.

```
// Add a car part to the assembly line
public void addCarPart(CarPart carPart) throws InterruptedException {
    try {
        emptyCount.acquire();
        mutex.acquire();
    } catch (InterruptedException e2) {
        e2.printStackTrace();
    }
    carParts.put(carPart);
    mutex.release();
    fillCount.release();
}
```

The consumer process (human worker) would remove the car parts from the assembly line and install them on the car.

```
// Remove a car part from the assembly line
public CarPart removeCarPart() throws InterruptedException {
    try {
        fillCount.acquire();
        mutex.acquire();
    } catch (InterruptedException e2) {
        e2.printStackTrace();
    }
    CarPart carPart = carParts.poll(1, TimeUnit.SECONDS);
    if (carPart != null) {
        mutex.release();
        emptyCount.release();
    }
    return carPart;
}
```

A semaphore would be used to synchronize access to the assembly line, which acts as the common buffer in this case.

When the robot wants to add a car part to the assembly line, it would acquire the semaphore. This would prevent the human worker from accessing the assembly line until the robot has finished adding the car part.

Once the robot has finished adding the car part, it would release the semaphore, allowing the human worker to access the assembly line and remove the car part.

The human worker would then install the car part on the car and repeat the process until all car parts have been assembled.